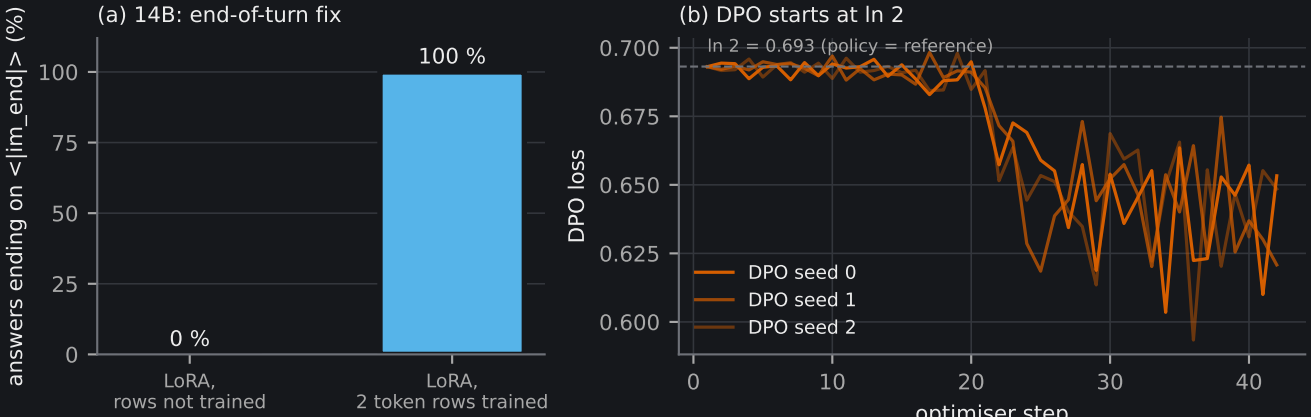
(a) 14B: end-of-turn fix
answers ending on <|im_end|> (%)
100
75
50
25
0
0 %
100 %
LoRA,
rows not trained
LoRA,
2 token rows trained
(b) DPO starts at ln 2
ln 2 = 0.693 (policy = reference)
DPO loss
0.700
0.675
0.650
0.625
0.600
DPO seed 0
DPO seed 1
DPO seed 2
0
10
20
30
40
optimiser step